

Prior authorization under pressure

Review Copilot
for Bellcourt Health Administrators



A review copilot that **picks the governing rule in code**, checks every criterion against the clinical record with a quoted citation, and leaves **every decision to a person**.

SUBMISSION

Presented by	Lohith Devaramane
Programme	FDE Academy, IIT Roorkee cohort
Evaluator	Sumit Shukla, FDE Academy
Date	7 October 2026
Format	8 minutes, then questions

WHAT IS IN THE BOX

- Working app: provider portal, intake and reviewer workspace, physician decisions, status desk, decision letters by PDF and email
- Evaluation against Bellcourt's own 120 audited cases, 30 live cases and 24 hostile inputs
- Documentation, architecture, design note, evidence pack and this deck



Riverbend is penalising Bellcourt for decisions that are late and wrong. Leadership says it is a staffing problem.

90.1%

Medicare Advantage decisions on time
Contract standard 97%. Corrective Action
Plan open since July 2026.

56.7%

Decisions correct in Bellcourt's own
QA audit
68 of 120 cases. Contract standard 95%.

56.2%

Riverbend denials overturned on
appeal
Root-cause review is triggered at 30%.
Wrong-version denials are overturned 91% of
the time.

\$1.9M

Service credits paid to Riverbend this
year
About a fifth of operating profit. Zero before
2026.

WHO BELLCOURT IS

A third-party administrator in Nashville. It runs self-funded health plans for **38 employers** (453,000 people) and, since 2023, decides prior authorization requests for **Riverbend Health Plan** (153,700 members). About **310,000 requests a year**, 46% arriving by fax.

WHAT IS AT STAKE

Riverbend's contract ends **31 December 2026** and may go out to bid this month. Harlan, the largest employer (about \$8.1M a year), renews on 1 January with satisfaction at 5 of 10. Operating margin has fallen from **14% to 8%** in three years.

THE CLAIM I HAD TO TEST

"It's a staffing problem, full stop." The VP of Operations wants 20 more nurses, about \$2.2M a year. The CEO is not convinced and wants decisions that are *"faster and more defensible, not just faster"*.

Four cuts of the data. Each one is a number, not an opinion.

CUT 1 · IS IT STAFFING?

$r = -0.05$

Nurse headcount against the on-time rate, week by week for 90 weeks.

Timeliness was already 87% when staffing was at 98% of budget. Headcount fell from 57.7 to 50.6 FTE and the on-time rate did not move with it.

No. Headcount does not move the number.

CUT 2 · WHY LATE?

64 of 64

missed Medicare deadlines arrived incomplete.

Clock rebuilt from fax receipt, as the contract counts it, not from keying 17 hours later. Complete requests were on time 99.8% of the time. In late cases 122 of 199 hours were spent waiting for missing information.

Lateness is an intake problem, not a review problem.

CUT 3 · WHY WRONG?

39 of 52

audit errors are about which rule was applied.

Sorted by the auditor's own label: memo overrode policy 18, retired policy version 14, client plan rule missed 7. Clinical misreads 10, information not requested 3.

The rulebook, not the nurse.

CUT 4 · WHAT DOES IT COST?

14.2 of 38

minutes per case spent finding the rule.

37% of nurse time, about 18.9 of 50.6 FTE, almost the 20 nurses the VP wants to hire. The senior nurse: "the hard part is figuring out which rules apply."

Fix the rule and the capacity appears.

Method: every claim tied to a file, every file reproduced by one script (evidence_analysis.py). Nothing on this slide comes from an interview alone; the interviews confirmed what the files said.

Two root causes, one thread: nobody can say which rule applies, and the clock starts late.

ROOT CAUSE 1 · THE WRONG RULEBOOK

34%

of 2026 denials cite a policy version that was already retired. Appeals against them succeed 91% of the time (30 of 33).

72

lumbar MRI requests denied under a memo that kept a 6-week rule after MP-101 v2 moved to 4 weeks. GOV-01 says a memo cannot change criteria.

8

Kestrel bariatric denials after an amendment that covered the surgery from January. The email went to claims configuration, never to UM.

20

glucose-monitor denials for Medicare members under the stricter Bellcourt rule, when the addendum says Medicare criteria apply. All 11 appeals overturned.

ROOT CAUSE 2 · THE LATE CLOCK

17 h

from fax arrival to keying. Riverbend's clock starts when the fax arrives; PACE starts it when someone types it in.

122 of 199 h

in late cases spent waiting for missing information. Nurse review took about 40 hours whether the case was late or not.

66 of 69


late decisions across all clients arrived incomplete; 64 of 64 for Medicare Advantage. Nobody tells the provider until a person opens the image.

1 in 3

faxes arrive missing something, says the intake coordinator: usually the member ID or the clinical notes. The provider finds out days later, by phone.

So what: fix the rule and the clock. Code can pick the governing document by client, service and date of service; code can stamp receipt and list what is missing in the first minute. Hiring nurses addresses neither.

One request, from the provider's desk to the written decision.



Provider portal: coverage check before filing, then submit

Nurse workspace: governing rule, criteria with quotes, decision


Decision letter from the human decision, PDF and email

- 1 **Provider** Coverage check for Harlan bariatric surgery: **not covered**, SPD-HARLAN \$6.4, before anything is filed.
- 2 **Provider** Submit the Juniper spinal-stimulator request. Receipt, request number and the clock at once; plan pre-check flags the second-opinion rule.
- 3 **Intake** Queue sorted by hours left. Open the new request, **run the review**: Need information, citing SPD-JUNIPER \$5.6. Pend it; the notice is created.
- 4 **Nurse** Case 8100, lumbar MRI: the memo and PACE say 6 weeks; the copilot picked MP-101 v2 by date of service and lists v1 and the memo as excluded. Approve.
- 5 **Physician** Case 8105 routed for Arizona sign-off. Deny with the stated reason; the letter carries criteria, citation and appeal rights.

Budget: 3 minutes 30 seconds. Steps 1 to 4 are the story; step 5 is optional. If the clock shows 7:00 and the demo is still running, stop and go to evaluation, the slide the evaluator asked for by name.




Six elements. Code picks the rule, the model checks the evidence, a person decides.




- **Deterministic where it matters.** Which document governs is never a model judgement; it is a lookup over effective dates and rank, and it is unit-tested.
- **Agentic RAG, bounded.** The model decides what to read and how to weigh evidence, inside a document set it cannot widen. Retrieval is filtered by the governing set, so a retired version cannot surface.
- **Every output is checked by code** before a person sees it, and every run stores inputs, documents, trace, output and verifier report.
- **Roles enforced on the server:** provider, intake, nurse, physician, member-services agent, auditor. Reviewers cannot raise cases; providers cannot see the review.
- **Portable:** one environment variable switches the model provider; the vector index and store sit behind one interface each.
- **Runs beside PACE,** read-only. Nothing is written back; PACE stays the system of record until its 2028 replacement.

Something for each person in the picture, and nothing that takes a decision away from them.




Provider portal

Coverage check before filing: eligibility, plan exclusions and conditions, the policy version that will apply, the documents to send. Submit by fax upload or form. Receipt and request number at once. Status and written notices, no phone call.




Status desk

For member-services agents. Request number and date of birth must both match before anything is disclosed. Status, due date and missing items only; never the clinical review. Every lookup audited.




Priority queue

Sorted by hours left on a clock that starts at receipt, on any channel. Overdue, urgent, due soon. Urgency conflicts between cover sheet and form are flagged and resolved to the stricter clock.




Review copilot

Governing documents with the reason each applies, and the ones deliberately not used. Each criterion met, not met or not documented with an exact quote. One click opens the source section.



Human decision and letters

Approve, pend, route, or physician-only deny. Overrides need a reason. The decision letter states the specific reason, the provision relied on and appeal rights; PDF download and email.




Governance page and audit log


Eleven live controls you can click to watch the server refuse, redact, flag or reject. Per-run trace: plan, reads, refusals, verifier, second check, tokens, latency.

Every constraint is enforced in code, and the app lets you watch it happen.

Constraint	How it is enforced	Proof on the Governance page
No automated denial, delay or modification (addendum §5)	The outcome schema has no deny value; the verifier rejects anything else. The tool drafts; it never sends.	Submit DENY: verifier rejection V1
Only physicians deny; Arizona MA needs an AZ-licensed director (§4)	Decision permissions by role on the server; state flag AZ_MD_SIGNOFF_REQUIRED routes the case and marks the letter.	Dry-run deny as a nurse: 403
Hierarchy of authority (GOV-01); memos never change criteria	Rule resolver by rank and effective date; read tool refuses non-governing documents; memos excluded with a reason.	Read MP-101 v1 for a 2026 date: REFUSED
Segregation of duties	Intake and providers raise requests; nurses and physicians decide; auditors read. Enforced on every request, not only in the interface.	Role matrix and live attempt
PHI minimisation	Name, date of birth, record number and phone stripped before the review model; only age passed. Fax images served only to signed-in users.	Before and after on a sample note
Untrusted provider content	Text that tries to direct the reviewer is flagged, shown to the human, and barred as evidence by the verifier.	Planted 'pre-approved' sentence flagged
Authentication and sessions	Salted PBKDF2 hashes, 200,000 rounds; HMAC-signed 8-hour sessions; constant-time checks.	Tampered and forged tokens rejected
Audit on request (addendum §5)	Every action with user, role and time; each review stores inputs, governing set, trace, output and verifier report.	Audit log and per-run trace



The Governance and security page. Each row above is a button; the server answers live.



Tested against the auditor's answers, hostile inputs and time travel, with every run traced.

117/120

agree with the QA auditor
Human reviewers on the same cases: 68/120

118/120

governing document and version correct
Document ID and version exact, section by prefix

0

wrongful approvals
No denial is possible by design

23/24

adversarial, incomplete and wrong-input tests
10 planted instructions all ignored; 7 time-travel cases right

29/30

open queue against my answer key
13 image-only faxes read, not guessed

WHERE THE HUMANS WENT WRONG, THE COPILOT WAS RIGHT

Memo override policy	18 of 18
Outdated policy version	14 of 14
Criteria misread	10 of 10
Client rule missed	6 of 7
Missing info not requested	2 of 3
No human error	67 of 68

WHAT EACH NUMBER MEASURES

Accuracy	Outcome equals the auditor's qa_correct_decision on 120 cases
Safety	Wrongful approvals and tool denials, both zero
Robustness	Planted text, missing fields, wrong dates, retired versions
Consistency	Repeat runs: 27 to 29 of 30, every miss toward a person
Observability	Plan, reads, refusals, verifier, second check, tokens, latency per run
Cost and speed	6.5s and about a third of a cent per case

THE THREE MISSES, ALL TOWARD A PERSON

QA-26-036 SORREL · Arthroscopic Knee Surgery · auditor deny medical necessity, copilot pend for info
QA-26-082 HARLAN · Hyperbaric Oxygen Therapy · auditor pend for info, copilot deny medical necessity
QA-26-095 BRIGHT · Bariatric · auditor deny not covered, copilot pend for info

Limits I state openly: results vary slightly between runs; prompts were tuned while looking at these cases, so a fresh random audit sample is the next test; the open-case key is my own reading of the policies; 6 of 38 employer plan documents are in the pack.

About \$0.55M to build and run in year one, against \$1.9M of credits and a \$2.2M hiring plan.

WHAT IT COSTS (\$M)

Item	Year 1	Ongoing
90-day build: FDE, 2 engineers, half a clinical SME, security review (A)	0.30	
Model and fax reading: 310k cases × \$0.003, 143k faxes × \$0.004	0.01	0.01
Azure hosting, Postgres, logging, monitoring (A)	0.05	0.06
Support: half an engineer, half a policy librarian (A)	0.12	0.24
Contingency 15% (A)	0.07	
Total	0.55	0.30

0.5% of revenue. One tenth of the hiring plan. Payback from credits alone in about five months.

WHAT IT IS WORTH, PER YEAR

- \$1.9M


service credits avoided. The FY2026 line. Each point of timeliness below 97% costs 4% of Riverbend's quarterly fees, about \$0.27M a year per point.
- \$2.2M

hiring avoided (A: 20 nurses at \$110k loaded). Finding the rule is 14.2 of 38 minutes per case, 37% of nurse time, about 18.9 FTE. Half of that fills the 7.4 FTE vacancy gap.
- 12.5 FTE

call-centre capacity. 67% of 2,400 daily calls are status or fax-receipt; half deflected by the portal and status desk at 446 s a call (A: 50%).
- \$6.8M

Riverbend fees protected, plus Harlan (\$8.1M) and 13 other employers (\$42M in admin fees) renewing on 1 January. Not counted in the margin bars.

OPERATING MARGIN, SCENARIOS



Revenue held at the FY2026 forecast; SG&A flat; no growth counted. The CEO's three further delegated plans are upside.

The argument in one line: the credits alone are 21% of operating profit and 1.7 margin points; the copilot costs 0.3 points a year and removes the cause, while the hiring plan costs 2 points and removes nothing.

Three decisions.

This week, no technology

- Withdraw UM-MEMO-2025-19
- Tell UM Operations about the Kestrel amendment
- Re-review denials still open to appeal under the old versions
- Report Riverbend timeliness from receipt time

90 days, recommend-only, beside PACE

- Send the 60-day notice to Riverbend on day 1
- Registry of documents owned by the Clinical Policy Committee
- Shadow mode with senior nurses, then self-funded clients, then Riverbend
- Provider portal and status desk to cut status calls

Measure it, from receipt

- MA on-time from about 90% to 97% or better
- QA accuracy from 56.7% to 95% or better
- Wrong-version citations to zero; overturns below 30%
- Stop rule: pause if accuracy on sampled cases falls under 90%

Fix the rule and the clock. Staffing follows.

Thank you. Questions?

